

Babar Hashmi

AI Engineer - Machine Learning Specialist

ibabarhashmi@yahoo.com | linkedin.com/in/hashmi-babar | github.com/ibabarhashmi
Bangalore, India | Open to Relocation

Professional Summary

AI engineer with 3 years of applied ML, most of it as the first AI hire at QuillAudits (now QuillAI Network). I started with no team and no existing pipeline, and built the models, agent workflows and developer SDK that the company now ships. Day to day that means fine-tuning LLMs for detection tasks, designing agent loops, and writing the docs and tests that keep results honest.

Professional Experience

Founding AI Engineer (MTS, Research Engineering) - QuillAI Network (QuillAudits) | Jul 2024 - Present

- > Shipped production deep-learning models for smart-contract vulnerability detection, hardened by red-teaming - catching what automated analysis misses.
- > First AI hire: set up the research-to-production process, then designed agent workflows with MCP tool-calling on open agent frameworks and shipped a developer SDK used across 3 business units.
- > Onboarded early users and debugged their deployments until things actually worked.
- > Fine-tuned LLMs for real detection tasks; rebuilt ML pipelines as the problem got clearer.

AI/ML Researcher - QuillAudits | Sep 2023 - Jun 2024

- > Built and launched QuillShield, an AI smart-contract auditing tool, as its lead AI engineer.
- > Ran statistical modelling across 100K+ smart contracts for vulnerability research.
- > Turned research into tools, including due-diligence analysis for market decisions.
- > Delivered with cross-functional Agile teams; wrote the internal docs that kept the work honest.

Key Projects

Financial Strategy Risk Scoring Engine (organon) - TypeScript

- > Deterministic scorecard over 7 risk axes from point-in-time market data; groundedness-gated LLM; one-command verification pipeline.

Agentic-Auditor - AI Smart Contract Auditor (TypeScript, Vite)

- > User-friendly, on-the-go audit UX with AI-assisted reporting.

GoldenHour - First response for cyber fraud (TypeScript)

- > Rapid-response tooling for cyber-fraud incidents.

Guardrails - Policy layer for LLM agents (Python)

- > Policy and guardrail layer keeping tool-calling agent loops grounded and safe.

DesiRAG - Hindi voice RAG (Python)

- > Hybrid retrieval over MSMARCO-XI with guardrails. Built for HH Goa 2026.

ERC20 Attribute Predictor - ERC20 ML (Python, Solidity)

- > ML model extracting attributes of ERC20 contracts from Solidity source and bytecode.

Vuln-Detector - Smart-contract vulnerability detection (Python)

- > Smart-contract vulnerability detection experiments.

MEP Industrial Elements Detection (GroundingDINO, AWS)

- > Open-set detection of MEP elements with real-time AWS inference API.

Real-Time Pothole Detection (YOLOv8, Kafka, Streamlit, AWS)

- > YOLOv8 detector with Kafka streaming pipeline, Streamlit dashboard and REST APIs.

Technical Proficiencies

ML/AI: PyTorch, TensorFlow, Transformers, Open agent frameworks, Custom agent frameworks, LLM fine-tuning. CV: YOLOv8, GroundingDINO, OpenCV, VLMs. Cloud/MLOps: AWS, OCI, Docker, CI/CD, Model monitoring, Inference optimization, vLLM. Eng: Python, TypeScript, Bun, Solidity,

SQL/NoSQL, REST APIs, TDD, Agent harness design, SDK design, Eval and red-teaming. Data: Pandas, NumPy, Matplotlib, Plotly, Streamlit.

Certifications

- > Agentic AI Certified Associate - Oracle - Jul 2026.
- > Copado Certified AI - Copado - Jul 2026.
- > OCI Generative AI Certified Professional - Oracle - Jul 2024.
- > Prompt & AI Engineering Certification - MFT - Jun 2024.

Impact Snapshot

+25% audit accuracy - -60% manual review - 93% ERC20 predictive-check accuracy - 7-axis deterministic scorecard - 100K+ contracts analysed.

Domains of Expertise

> Blockchain / Web3 security; Construction tech (MEP); Public infrastructure CV; FinTech risk scoring; Multi-modal AI systems.